

Assignment1_yli130

September 24, 2023

1 Assignment 1: Neural Networks

Yanxi Li, yli130@kent.edu

1.0.1 The IMDB dataset

Loading the IMDB dataset

```
[1]: from tensorflow.keras.datasets import imdb
      (train_data, train_labels), (test_data, test_labels) = imdb.load_data(
          num_words=10000)
```

```
[2]: import matplotlib.pyplot as plt # for figure plot
```

```
[3]: len(train_data) # comments in training data set
```

```
[3]: 25000
```

```
[4]: len(train_data[1]) # check the length of first comment in training data
```

```
[4]: 189
```

```
[5]: max_len = max([len(i) for i in train_data]) #check max length in each list in
      ↪ the training data
      max_len
```

```
[5]: 2494
```

```
[6]: max([max(sequence) for sequence in train_data]) # check the maximum number in
      ↪ the list
```

```
[6]: 9999
```

1.0.2 Preparing the data

Encoding the integer sequences via multi-hot encoding

```
[3]: import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
[8]: x_train[1] # result for the multi-hot encoding for index 1
```

```
[8]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
[9]: x_train.shape, x_test.shape # show both train_test data set shape
```

```
[9]: ((25000, 10000), (25000, 10000))
```

```
[10]: len(x_train[0]) # check the length for the first sample in training data set
```

```
[10]: 10000
```

```
[4]: y_train = np.asarray(train_labels).astype("float32") # convert both train &
      ↪ test labels into floats
y_test = np.asarray(test_labels).astype("float32")
```

```
[12]: type(x_train), type(y_train) #check data format in training data set
```

```
[12]: (numpy.ndarray, numpy.ndarray)
```

```
[13]: y_train.shape, y_test.shape # check train and test label shape
```

```
[13]: ((25000,), (25000,))
```

1.0.3 Original model

Build the model

```
[5]: # Model definition
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

```

# Compile the model
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

# Setting the validation set
x_val = x_train[:10000]           # validation for the first 10000 data in
    ↳ training set
partial_x_train = x_train[10000:] # training data for the rest of data
    ↳ which is 25000-10000=15000
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

# Train the model
history = model.fit(partial_x_train,
                   partial_y_train,
                   epochs=20,      # run the model for 20 epoch and each epoch
    ↳ for 512 batch
                   batch_size=512,
                   validation_data=(x_val, y_val))

```

```

Epoch 1/20
30/30 [=====] - 2s 41ms/step - loss: 0.5497 - accuracy:
0.7778 - val_loss: 0.4259 - val_accuracy: 0.8672
Epoch 2/20
30/30 [=====] - 1s 24ms/step - loss: 0.3458 - accuracy:
0.8992 - val_loss: 0.3439 - val_accuracy: 0.8668
Epoch 3/20
30/30 [=====] - 1s 22ms/step - loss: 0.2524 - accuracy:
0.9224 - val_loss: 0.2913 - val_accuracy: 0.8895
Epoch 4/20
30/30 [=====] - 1s 24ms/step - loss: 0.1966 - accuracy:
0.9373 - val_loss: 0.2869 - val_accuracy: 0.8855
Epoch 5/20
30/30 [=====] - 1s 24ms/step - loss: 0.1635 - accuracy:
0.9473 - val_loss: 0.2824 - val_accuracy: 0.8880
Epoch 6/20
30/30 [=====] - 1s 24ms/step - loss: 0.1344 - accuracy:
0.9595 - val_loss: 0.2829 - val_accuracy: 0.8862
Epoch 7/20
30/30 [=====] - 1s 21ms/step - loss: 0.1156 - accuracy:
0.9658 - val_loss: 0.2958 - val_accuracy: 0.8837
Epoch 8/20
30/30 [=====] - 1s 22ms/step - loss: 0.0944 - accuracy:
0.9740 - val_loss: 0.3133 - val_accuracy: 0.8818
Epoch 9/20
30/30 [=====] - 1s 22ms/step - loss: 0.0794 - accuracy:

```

```

0.9789 - val_loss: 0.3397 - val_accuracy: 0.8780
Epoch 10/20
30/30 [=====] - 1s 21ms/step - loss: 0.0659 - accuracy:
0.9829 - val_loss: 0.3759 - val_accuracy: 0.8761
Epoch 11/20
30/30 [=====] - 1s 24ms/step - loss: 0.0564 - accuracy:
0.9859 - val_loss: 0.3818 - val_accuracy: 0.8774
Epoch 12/20
30/30 [=====] - 1s 24ms/step - loss: 0.0459 - accuracy:
0.9901 - val_loss: 0.4082 - val_accuracy: 0.8747
Epoch 13/20
30/30 [=====] - 1s 24ms/step - loss: 0.0375 - accuracy:
0.9919 - val_loss: 0.4532 - val_accuracy: 0.8691
Epoch 14/20
30/30 [=====] - 1s 21ms/step - loss: 0.0313 - accuracy:
0.9941 - val_loss: 0.4610 - val_accuracy: 0.8729
Epoch 15/20
30/30 [=====] - 1s 21ms/step - loss: 0.0251 - accuracy:
0.9956 - val_loss: 0.4997 - val_accuracy: 0.8704
Epoch 16/20
30/30 [=====] - 1s 24ms/step - loss: 0.0199 - accuracy:
0.9971 - val_loss: 0.5239 - val_accuracy: 0.8699
Epoch 17/20
30/30 [=====] - 1s 24ms/step - loss: 0.0149 - accuracy:
0.9985 - val_loss: 0.5921 - val_accuracy: 0.8632
Epoch 18/20
30/30 [=====] - 1s 20ms/step - loss: 0.0148 - accuracy:
0.9974 - val_loss: 0.5924 - val_accuracy: 0.8678
Epoch 19/20
30/30 [=====] - 1s 18ms/step - loss: 0.0103 - accuracy:
0.9989 - val_loss: 0.6206 - val_accuracy: 0.8677
Epoch 20/20
30/30 [=====] - 1s 18ms/step - loss: 0.0066 - accuracy:
0.9997 - val_loss: 0.6738 - val_accuracy: 0.8637

```

```

[6]: history_dict = history.history
     history_dict.keys() # show key names in the dictionary

```

```

[6]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```


Plotting the training and validation loss

```

[7]: history_dict = history.history
     loss_values = history_dict["loss"]
     val_loss_values = history_dict["val_loss"]
     epochs = range(1, len(loss_values) + 1)
     plt.plot(epochs, loss_values, "b", label="Training loss")
     plt.plot(epochs, val_loss_values, "r", label="Validation loss")

```

```
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



From the loss pattern in training & validation, the validation loss **starts to increase** around 4 epochs, the **original model choose the 4 epochs** to retrain the model using the testing data.

Retraining a model from scratch

```
[17]: model = keras.Sequential([
        layers.Dense(16, activation="relu"),
        layers.Dense(16, activation="relu"),
        layers.Dense(1, activation="sigmoid")
    ])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=4, batch_size=512) # change to 4 epochs
results = model.evaluate(x_test, y_test)
```

Epoch 1/4

49/49 [=====] - 2s 21ms/step - loss: 0.4621 - accuracy:

```

0.8190
Epoch 2/4
49/49 [=====] - 1s 17ms/step - loss: 0.2676 - accuracy:
0.9075
Epoch 3/4
49/49 [=====] - 1s 14ms/step - loss: 0.2040 - accuracy:
0.9261
Epoch 4/4
49/49 [=====] - 1s 13ms/step - loss: 0.1698 - accuracy:
0.9404
782/782 [=====] - 1s 573us/step - loss: 0.2900 -
accuracy: 0.8840

```

[21]: results

[21]: [0.29001015424728394, 0.8840000033378601]

The original model **test data prediction** result: **loss: 0.2926, accuracy: 0.8840**.

1.1 Question 1: Change to three hidden layers

```

[18]: # Build the model with 3 layers
model_3_layers = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_3_layers.compile(optimizer="rmsprop",
                       loss="binary_crossentropy",
                       metrics=["accuracy"])

# Train the model
history_3_layers = model_3_layers.fit(partial_x_train,
                                     partial_y_train,
                                     epochs=20,      # run the model for 20 epoch and each epoch
                                     ↪for 512 batch
                                     batch_size=512,
                                     validation_data=(x_val, y_val))

```

```

Epoch 1/20
30/30 [=====] - 2s 35ms/step - loss: 0.5298 - accuracy:
0.7732 - val_loss: 0.3966 - val_accuracy: 0.8613
Epoch 2/20
30/30 [=====] - 1s 23ms/step - loss: 0.3064 - accuracy:
0.9027 - val_loss: 0.3038 - val_accuracy: 0.8858

```

Epoch 3/20
30/30 [=====] - 1s 21ms/step - loss: 0.2173 - accuracy:
0.9287 - val_loss: 0.3364 - val_accuracy: 0.8625
Epoch 4/20
30/30 [=====] - 1s 21ms/step - loss: 0.1714 - accuracy:
0.9434 - val_loss: 0.2753 - val_accuracy: 0.8910
Epoch 5/20
30/30 [=====] - 1s 23ms/step - loss: 0.1416 - accuracy:
0.9526 - val_loss: 0.2970 - val_accuracy: 0.8845
Epoch 6/20
30/30 [=====] - 1s 21ms/step - loss: 0.1171 - accuracy:
0.9625 - val_loss: 0.3157 - val_accuracy: 0.8821
Epoch 7/20
30/30 [=====] - 1s 23ms/step - loss: 0.0951 - accuracy:
0.9706 - val_loss: 0.3732 - val_accuracy: 0.8660
Epoch 8/20
30/30 [=====] - 1s 24ms/step - loss: 0.0800 - accuracy:
0.9759 - val_loss: 0.3818 - val_accuracy: 0.8668
Epoch 9/20
30/30 [=====] - 1s 21ms/step - loss: 0.0692 - accuracy:
0.9797 - val_loss: 0.3752 - val_accuracy: 0.8744
Epoch 10/20
30/30 [=====] - 1s 27ms/step - loss: 0.0552 - accuracy:
0.9839 - val_loss: 0.4244 - val_accuracy: 0.8738
Epoch 11/20
30/30 [=====] - 1s 31ms/step - loss: 0.0447 - accuracy:
0.9881 - val_loss: 0.4395 - val_accuracy: 0.8667
Epoch 12/20
30/30 [=====] - 1s 28ms/step - loss: 0.0379 - accuracy:
0.9903 - val_loss: 0.5489 - val_accuracy: 0.8503
Epoch 13/20
30/30 [=====] - 1s 30ms/step - loss: 0.0310 - accuracy:
0.9923 - val_loss: 0.5070 - val_accuracy: 0.8730
Epoch 14/20
30/30 [=====] - 1s 24ms/step - loss: 0.0263 - accuracy:
0.9936 - val_loss: 0.5220 - val_accuracy: 0.8703
Epoch 15/20
30/30 [=====] - 1s 25ms/step - loss: 0.0219 - accuracy:
0.9945 - val_loss: 0.5650 - val_accuracy: 0.8695
Epoch 16/20
30/30 [=====] - 1s 24ms/step - loss: 0.0118 - accuracy:
0.9987 - val_loss: 0.8162 - val_accuracy: 0.8320
Epoch 17/20
30/30 [=====] - 1s 24ms/step - loss: 0.0137 - accuracy:
0.9971 - val_loss: 0.6260 - val_accuracy: 0.8630
Epoch 18/20
30/30 [=====] - 1s 21ms/step - loss: 0.0160 - accuracy:
0.9956 - val_loss: 0.6589 - val_accuracy: 0.8660

Epoch 19/20


30/30 [=====] - 1s 21ms/step - loss: 0.0049 - accuracy: 0.9998 - val_loss: 0.7011 - val_accuracy: 0.8653

Epoch 20/20

30/30 [=====] - 1s 20ms/step - loss: 0.0116 - accuracy: 0.9970 - val_loss: 0.7284 - val_accuracy: 0.8638

Train & Validation loss

```
[19]: history_dict_3layers = history_3_layers.history
loss_values = history_dict_3layers["loss"]
val_loss_values = history_dict_3layers["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "b", label="Training loss")
plt.plot(epochs, val_loss_values, "r", label="Validation loss")
plt.title("Training and validation loss with 3 hidden layers")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



The 3 hiddel layer model still need to **choose 4 epochs** to run the test data.

Test accuracy


```
[20]: model_3_layers = keras.Sequential([
        layers.Dense(16, activation="relu"),
        layers.Dense(16, activation="relu"),
        layers.Dense(16, activation="relu"),
        layers.Dense(1, activation="sigmoid")
    ])
model_3_layers.compile(optimizer="rmsprop",
                       loss="binary_crossentropy",
                       metrics=["accuracy"])
model_3_layers.fit(x_train, y_train, epochs=4, batch_size=512) # change to 4
↳ epochs
results_3_layers = model_3_layers.evaluate(x_test, y_test)
```

```
Epoch 1/4
49/49 [=====] - 2s 19ms/step - loss: 0.4586 - accuracy:
0.8158 0s - loss: 0.5132 - accuracy
Epoch 2/4
49/49 [=====] - 1s 15ms/step - loss: 0.2501 - accuracy:
0.9092
Epoch 3/4
49/49 [=====] - 1s 13ms/step - loss: 0.1936 - accuracy:
0.9316
Epoch 4/4
49/49 [=====] - 1s 14ms/step - loss: 0.1618 - accuracy:
0.9426
782/782 [=====] - 1s 615us/step - loss: 0.3106 -
accuracy: 0.8806
```

```
[22]: results_3_layers
```

```
[22]: [0.31061962246894836, 0.8805999755859375]
```

The 3 hidden layers model test data prediction result: **loss: 0.3108, accuracy: 0.8805**.

Validation accuracy & loss comparison

```
[23]: #plt.rcParams['figure.dpi'] = 300
history_dict_3layer = history_3_layers.history
val_acc_3_layers = history_dict_3layer['val_accuracy']
val_acc = history_dict['val_accuracy']

val_loss_3_layers = history_dict_3layer['val_loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(val_acc)+1)

plt.figure(figsize=(10,4)) # set the figure size
```

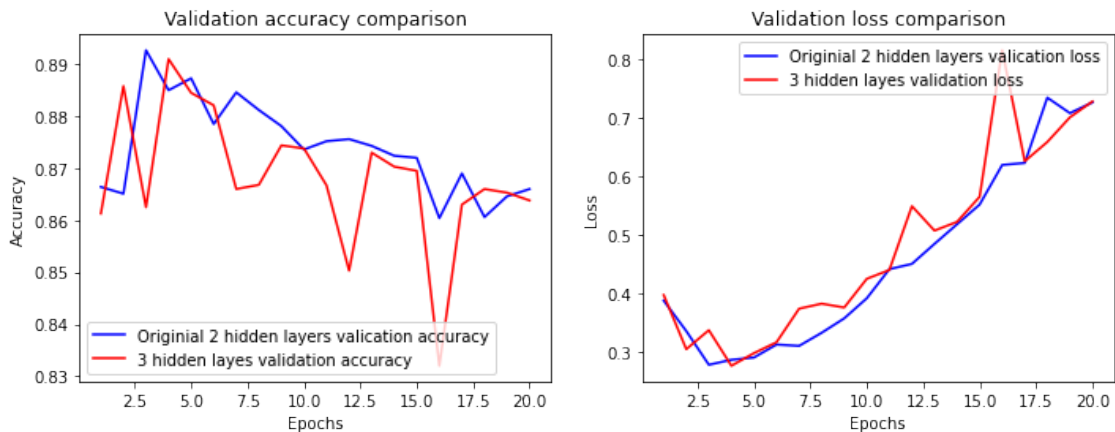
```

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="Original 2 hidden layers validation_
↪accuracy")
plt.plot(epochs, val_acc_3_layers, "r", label="3 hidden layes validation_
↪accuracy")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="Original 2 hidden layers valication_
↪loss")
plt.plot(epochs, val_loss_3_layers, "r", label="3 hidden layes validation loss")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()

```



1.1.1 Summary for Question 1:

Similarly, 4 epoch is also appropriate for the 3 hidden layers to **minimize the loss function as well as avoid overfit**. Validation accuracy in both models(2&3 hidden layers) are **close to each other**. As for the 4 epoch, 3 hidden layers **accuracy is even lower** than the original 2 hidden layers

Model with 3 hidden layers test accuracy is 0.8805, while the original model with 2 hidden layers test accuracy is 0.8840. Increasing the number of hidden layers from 2 to 3 even **lower the test accuracy**.

1.2 Question 2: 32 units with 2 hidden layers

```
[8]: # Build the model with 32units and 2 layers
model_32units = keras.Sequential([
    layers.Dense(32, activation="relu"),
    layers.Dense(32, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_32units.compile(optimizer="rmsprop",
                      loss="binary_crossentropy",
                      metrics=["accuracy"])

# Train the model
history_32units = model_32units.fit(partial_x_train,
                                    partial_y_train,
                                    epochs=20,      # run the model for 20 epochs and each epoch
                                    ↪for 512 batch
                                    batch_size=512,
                                    validation_data=(x_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 3s 42ms/step - loss: 0.4974 - accuracy:
0.7850 - val_loss: 0.3971 - val_accuracy: 0.8410
Epoch 2/20
30/30 [=====] - 1s 28ms/step - loss: 0.2802 - accuracy:
0.9047 - val_loss: 0.2906 - val_accuracy: 0.8891
Epoch 3/20
30/30 [=====] - 1s 25ms/step - loss: 0.2055 - accuracy:
0.9273 - val_loss: 0.2813 - val_accuracy: 0.8880
Epoch 4/20
30/30 [=====] - 1s 27ms/step - loss: 0.1597 - accuracy:
0.9462 - val_loss: 0.2829 - val_accuracy: 0.8889
Epoch 5/20
30/30 [=====] - 1s 27ms/step - loss: 0.1283 - accuracy:
0.9569 - val_loss: 0.2932 - val_accuracy: 0.8859
Epoch 6/20
30/30 [=====] - 1s 28ms/step - loss: 0.1030 - accuracy:
0.9672 - val_loss: 0.3135 - val_accuracy: 0.8852
Epoch 7/20
30/30 [=====] - 1s 34ms/step - loss: 0.0832 - accuracy:
0.9737 - val_loss: 0.3742 - val_accuracy: 0.8754
Epoch 8/20
30/30 [=====] - 1s 35ms/step - loss: 0.0671 - accuracy:
0.9798 - val_loss: 0.3831 - val_accuracy: 0.8781
Epoch 9/20
30/30 [=====] - 1s 25ms/step - loss: 0.0480 - accuracy:
```

```

0.9883 - val_loss: 0.4038 - val_accuracy: 0.8741
Epoch 10/20
30/30 [=====] - 1s 25ms/step - loss: 0.0427 - accuracy:
0.9879 - val_loss: 0.4442 - val_accuracy: 0.8750
Epoch 11/20
30/30 [=====] - 1s 25ms/step - loss: 0.0280 - accuracy:
0.9943 - val_loss: 0.6271 - val_accuracy: 0.8494
Epoch 12/20
30/30 [=====] - 1s 25ms/step - loss: 0.0289 - accuracy:
0.9913 - val_loss: 0.4948 - val_accuracy: 0.8737
Epoch 13/20
30/30 [=====] - 1s 27ms/step - loss: 0.0165 - accuracy:
0.9970 - val_loss: 0.6834 - val_accuracy: 0.8439
Epoch 14/20
30/30 [=====] - 1s 25ms/step - loss: 0.0097 - accuracy:
0.9993 - val_loss: 0.5662 - val_accuracy: 0.8711
Epoch 15/20
30/30 [=====] - 1s 27ms/step - loss: 0.0165 - accuracy:
0.9948 - val_loss: 0.6008 - val_accuracy: 0.8702
Epoch 16/20
30/30 [=====] - 1s 27ms/step - loss: 0.0044 - accuracy:
0.9998 - val_loss: 0.6510 - val_accuracy: 0.8647
Epoch 17/20
30/30 [=====] - 1s 24ms/step - loss: 0.0099 - accuracy:
0.9972 - val_loss: 0.6760 - val_accuracy: 0.8664
Epoch 18/20
30/30 [=====] - 1s 31ms/step - loss: 0.0023 - accuracy:
0.9999 - val_loss: 0.7145 - val_accuracy: 0.8664
Epoch 19/20
30/30 [=====] - 1s 28ms/step - loss: 0.0091 - accuracy:
0.9967 - val_loss: 0.7447 - val_accuracy: 0.8657
Epoch 20/20
30/30 [=====] - 1s 28ms/step - loss: 0.0012 - accuracy:
0.9999 - val_loss: 0.7755 - val_accuracy: 0.8656

```

Train & Validation loss

```

[9]: history_dict_32units = history_32units.history


loss_values = history_dict_32units['loss'][:10]  # select 10 epoch could be
→easier to detect the suitable epoch.
val_loss_values = history_dict_32units['val_loss'][:10]

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, "b", label="Training loss")
plt.plot(epochs, val_loss_values, "r", label="Validation loss")
plt.title("Training and validation loss with 32 units & 2 layers")

```

```
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



For the 32 units with 2 layers, I choose **3 epochs** to run the test data.

Test accuracy

```
[10]: model_32units = keras.Sequential([
        layers.Dense(32, activation="relu"),
        layers.Dense(32, activation="relu"),
        layers.Dense(1, activation="sigmoid")
    ])
model_32units.compile(optimizer="rmsprop",
                      loss="binary_crossentropy",
                      metrics=["accuracy"])
model_32units.fit(x_train, y_train, epochs=3, batch_size=512) # change to 3
↳ epochs
results_32units = model_32units.evaluate(x_test, y_test)
```

Epoch 1/3

49/49 [=====] - 2s 25ms/step - loss: 0.4352 - accuracy: 0.8206

```
Epoch 2/3
49/49 [=====] - 1s 17ms/step - loss: 0.2392 - accuracy:
0.9138
Epoch 3/3
49/49 [=====] - 1s 17ms/step - loss: 0.1929 - accuracy:
0.9289
782/782 [=====] - 1s 706us/step - loss: 0.2894 -
accuracy: 0.8851
```

```
[11]: results_32units
```

```
[11]: [0.2894221544265747, 0.8850799798965454]
```

The 32 units with 2 layers model test data prediction result: **loss: 0.2894, accuracy: 0.8850.**

Validation accuracy & loss comparison

```
[12]: val_acc_32units = history_dict_32units['val_accuracy']
val_acc = history_dict['val_accuracy']

val_loss_32units = history_dict_32units['val_loss']
val_loss = history_dict['val_loss']


epochs = range(1, len(val_acc_32units)+1)

plt.figure(figsize=(10,4)) # set the figure size

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="Original 16 untis valication accuracy")
plt.plot(epochs, val_acc_32units, "r", label="32 units validation accuracy")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="Original 16 units valication loss")
plt.plot(epochs, val_loss_32units, "r", label="32 units validation loss")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()
```



1.2.1 Summary for Question 2:

3 epoch is appropriate for the 32 units with 2 layers to **minimize the loss function as well as avoid overfit**. **Validation accuracy in 32 units is lower** than the original model, and from the validation loss, **32units begins overfit ealier** than the original 16units model.

Model with **32 units test accuracy is 0.8850**, while the original model with **16 units test accuracy is 0.8840**. Changing the number of units from 16 to 32 **increase the test accuracy**.

1.3 Question 3: Using MSE loss function

```
[13]: # Build the model with 16units and 2 layers
model_mse = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_mse.compile(optimizer="rmsprop",
                  loss="MSE",
                  metrics=["accuracy"])

# Train the model
history_mse = model_mse.fit(partial_x_train,
                           partial_y_train,
                           epochs=20,      # run the model for 20 epochs and each epoch
                           ↪for 512 batch
                           batch_size=512,
                           validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 2s 38ms/step - loss: 0.1650 - accuracy:

0.7911 - val_loss: 0.1150 - val_accuracy: 0.8745
Epoch 2/20
30/30 [=====] - 1s 21ms/step - loss: 0.0873 - accuracy:
0.9063 - val_loss: 0.0959 - val_accuracy: 0.8818
Epoch 3/20
30/30 [=====] - 1s 24ms/step - loss: 0.0629 - accuracy:
0.9313 - val_loss: 0.0852 - val_accuracy: 0.8878
Epoch 4/20
30/30 [=====] - 1s 28ms/step - loss: 0.0497 - accuracy:
0.9457 - val_loss: 0.0877 - val_accuracy: 0.8824
Epoch 5/20
30/30 [=====] - 1s 35ms/step - loss: 0.0396 - accuracy:
0.9599 - val_loss: 0.0865 - val_accuracy: 0.8816
Epoch 6/20
30/30 [=====] - 1s 24ms/step - loss: 0.0330 - accuracy:
0.9679 - val_loss: 0.0866 - val_accuracy: 0.8817
Epoch 7/20
30/30 [=====] - 1s 25ms/step - loss: 0.0276 - accuracy:
0.9729 - val_loss: 0.0879 - val_accuracy: 0.8818
Epoch 8/20
30/30 [=====] - 1s 23ms/step - loss: 0.0227 - accuracy:
0.9791 - val_loss: 0.0885 - val_accuracy: 0.8793
Epoch 9/20
30/30 [=====] - 1s 21ms/step - loss: 0.0177 - accuracy:
0.9846 - val_loss: 0.0915 - val_accuracy: 0.8785
Epoch 10/20
30/30 [=====] - 1s 24ms/step - loss: 0.0160 - accuracy:
0.9863 - val_loss: 0.0937 - val_accuracy: 0.8752
Epoch 11/20
30/30 [=====] - 1s 24ms/step - loss: 0.0130 - accuracy:
0.9894 - val_loss: 0.0958 - val_accuracy: 0.8736
Epoch 12/20
30/30 [=====] - 1s 21ms/step - loss: 0.0105 - accuracy:
0.9918 - val_loss: 0.0996 - val_accuracy: 0.8695
Epoch 13/20
30/30 [=====] - 1s 21ms/step - loss: 0.0093 - accuracy:
0.9922 - val_loss: 0.1002 - val_accuracy: 0.8718
Epoch 14/20
30/30 [=====] - 1s 23ms/step - loss: 0.0077 - accuracy:
0.9949 - val_loss: 0.1030 - val_accuracy: 0.8678
Epoch 15/20
30/30 [=====] - 1s 21ms/step - loss: 0.0067 - accuracy:
0.9943 - val_loss: 0.1039 - val_accuracy: 0.8702
Epoch 16/20
30/30 [=====] - 1s 24ms/step - loss: 0.0050 - accuracy:
0.9967 - val_loss: 0.1053 - val_accuracy: 0.8674
Epoch 17/20
30/30 [=====] - 1s 21ms/step - loss: 0.0060 - accuracy:


```

0.9948 - val_loss: 0.1075 - val_accuracy: 0.8670
Epoch 18/20
30/30 [=====] - 1s 20ms/step - loss: 0.0033 - accuracy:
0.9975 - val_loss: 0.1091 - val_accuracy: 0.8651
Epoch 19/20
30/30 [=====] - 1s 17ms/step - loss: 0.0054 - accuracy:
0.9947 - val_loss: 0.1096 - val_accuracy: 0.8652
Epoch 20/20
30/30 [=====] - 1s 20ms/step - loss: 0.0027 - accuracy:
0.9976 - val_loss: 0.1115 - val_accuracy: 0.8627

```

Train & Validation loss

```


[14]: history_dict_mse = history_mse.history

loss_values = history_dict_mse['loss'][:10]  # select 10 epoch could be easier
↳to detect the suitable epoch.
val_loss_values = history_dict_mse['val_loss'][:10]

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, "b", label="Training loss")
plt.plot(epochs, val_loss_values, "r", label="Validation loss")
plt.title("Training and validation loss with MSE loss function")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```



The MSE loss function model still need to choose **4 epochs** to run the test data. But we **cannot tell any increase from the validation loss** curve.

Test accuracy

```
[15]: model_mse = keras.Sequential([
        layers.Dense(16, activation="relu"),
        layers.Dense(16, activation="relu"),
        layers.Dense(1, activation="sigmoid")
    ])
model_mse.compile(optimizer="rmsprop",
                  loss="MSE",
                  metrics=["accuracy"])
model_mse.fit(x_train, y_train, epochs=4, batch_size=512) # change to 3 epochs
results_mse = model_mse.evaluate(x_test, y_test)
```

Epoch 1/4

49/49 [=====] - 2s 17ms/step - loss: 0.1477 - accuracy: 0.8126

Epoch 2/4

49/49 [=====] - 1s 19ms/step - loss: 0.0777 - accuracy: 0.9104

Epoch 3/4

49/49 [=====] - 1s 15ms/step - loss: 0.0586 - accuracy: 0.9310

```
Epoch 4/4
49/49 [=====] - 1s 13ms/step - loss: 0.0472 - accuracy:
0.9456
782/782 [=====] - 1s 679us/step - loss: 0.0889 -
accuracy: 0.8804
```

```
[16]: results_mse
```

```
[16]: [0.08892610669136047, 0.8804000020027161]
```

The MSE loss function model test data prediction result: **loss: 0.089, accuracy: 0.8804.**

Validation accuracy & loss comparison

```
[17]: val_acc_mse = history_dict_mse['val_accuracy']
val_acc = history_dict['val_accuracy']

val_loss_mse = history_dict_mse['val_loss']
val_loss = history_dict['val_loss']


epochs = range(1, len(val_acc_mse)+1)

plt.figure(figsize=(10,4)) # set the figure size

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="binary_crossentropy")
plt.plot(epochs, val_acc_mse, "r", label="MSE loss function")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="binary_crossentropy")
plt.plot(epochs, val_loss_mse, "r", label="MSE loss function")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()
```



1.3.1 Summary for Question 3:

Although the loss function MSE doesnot change much in test data loss and accuracy as well as validation data accuracy, the validation loss figure shows that the MSE loss doesnot change with the epochs increasing. MSE loss function should not be used in this condition.

Distinguish binary_crossentropy & MSE Loss function **binary_crossentropy** is a loss function that is often used for **binary classification** problems, where the output is a probability distribution over a set of classes. Cross entropy loss function penalizes the neural network for assigning low probabilities to the correct class and high probabilities to the wrong classes.

MSE, **mean squared error**, is often used for **regression** problems, where the output is a continuous value or a vector of values. MSE penalizes the neural network for deviating from the true output by squaring the difference between them.

1.4 Question 4: Using tanh activation

```
[8]: # Build the model with 16 units and 2 layers
model_tanh = keras.Sequential([
    layers.Dense(16, activation="tanh"),
    layers.Dense(16, activation="tanh"),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_tanh.compile(optimizer="rmsprop",
                   loss="binary_crossentropy",
                   metrics=["accuracy"])

# Train the model
history_tanh = model_tanh.fit(partial_x_train,
                             partial_y_train,
```

```

epochs=10,      # run the model for 10 epochs and each epoch
↳for 512 batch
batch_size=512,
validation_data=(x_val, y_val))

```

```

Epoch 1/10
30/30 [=====] - 2s 35ms/step - loss: 0.4712 - accuracy:
0.7972 - val_loss: 0.3415 - val_accuracy: 0.8754
Epoch 2/10
30/30 [=====] - 1s 21ms/step - loss: 0.2634 - accuracy:
0.9083 - val_loss: 0.2859 - val_accuracy: 0.8854
Epoch 3/10
30/30 [=====] - 1s 24ms/step - loss: 0.1854 - accuracy:
0.9377 - val_loss: 0.2685 - val_accuracy: 0.8904
Epoch 4/10
30/30 [=====] - 1s 21ms/step - loss: 0.1444 - accuracy:
0.9485 - val_loss: 0.2849 - val_accuracy: 0.8872
Epoch 5/10
30/30 [=====] - 1s 21ms/step - loss: 0.1129 - accuracy:
0.9617 - val_loss: 0.3089 - val_accuracy: 0.8836
Epoch 6/10
30/30 [=====] - 1s 21ms/step - loss: 0.0828 - accuracy:
0.9735 - val_loss: 0.4691 - val_accuracy: 0.8535
Epoch 7/10
30/30 [=====] - 1s 21ms/step - loss: 0.0661 - accuracy:
0.9786 - val_loss: 0.3976 - val_accuracy: 0.8730
Epoch 8/10
30/30 [=====] - 1s 20ms/step - loss: 0.0556 - accuracy:
0.9829 - val_loss: 0.4322 - val_accuracy: 0.8731
Epoch 9/10
30/30 [=====] - 1s 27ms/step - loss: 0.0454 - accuracy:
0.9861 - val_loss: 0.4628 - val_accuracy: 0.8730
Epoch 10/10
30/30 [=====] - 1s 24ms/step - loss: 0.0278 - accuracy:
0.9929 - val_loss: 0.5062 - val_accuracy: 0.8697

```

Train & validation loss

```

[9]: history_dict_tanh = history_tanh.history


loss_values = history_dict_tanh['loss'][:10] # select 10 epoch could be easier
↳to detect the suitable epoch.
val_loss_values = history_dict_tanh['val_loss'][:10]

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, "b", label="Training loss")
plt.plot(epochs, val_loss_values, "r", label="Validation loss")

```

```
plt.title("Training and validation with tanh activation")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



For the tanh activation model with 16 units and 2 layers, I choose **3 epochs** to run the test data.

Test accuracy

```
[10]: model_tanh = keras.Sequential([
        layers.Dense(16, activation="tanh"),
        layers.Dense(16, activation="tanh"),
        layers.Dense(1, activation="sigmoid")
    ])
model_tanh.compile(optimizer="rmsprop",
                    loss="binary_crossentropy",
                    metrics=["accuracy"])
model_tanh.fit(x_train, y_train, epochs=3, batch_size=512) # change to 3 epochs
results_tanh = model_tanh.evaluate(x_test, y_test)
```

Epoch 1/3

49/49 [=====] - 2s 21ms/step - loss: 0.4226 - accuracy: 0.8226

```
Epoch 2/3
49/49 [=====] - 1s 14ms/step - loss: 0.2289 - accuracy:
0.9153
Epoch 3/3
49/49 [=====] - 1s 13ms/step - loss: 0.1800 - accuracy:
0.9323 0s - loss: 0.1789 - accuracy - ETA: 0s - loss: 0.1813 - accuracy: 0.93
782/782 [=====] - 1s 620us/step - loss: 0.3003 -
accuracy: 0.8820
```

```
[11]: results_tanh
```

```
[11]: [0.3002675473690033, 0.8819599747657776]
```

The tanh activation model test data prediction result: **loss: 0.30626, accuracy: 0.8820.**

Validation accuracy & loss comparison

```
[13]: val_acc_tanh = history_dict_tanh['val_accuracy']
val_acc = history_dict['val_accuracy'][:10]

val_loss_tanh = history_dict_tanh['val_loss']
val_loss = history_dict['val_loss'][:10]

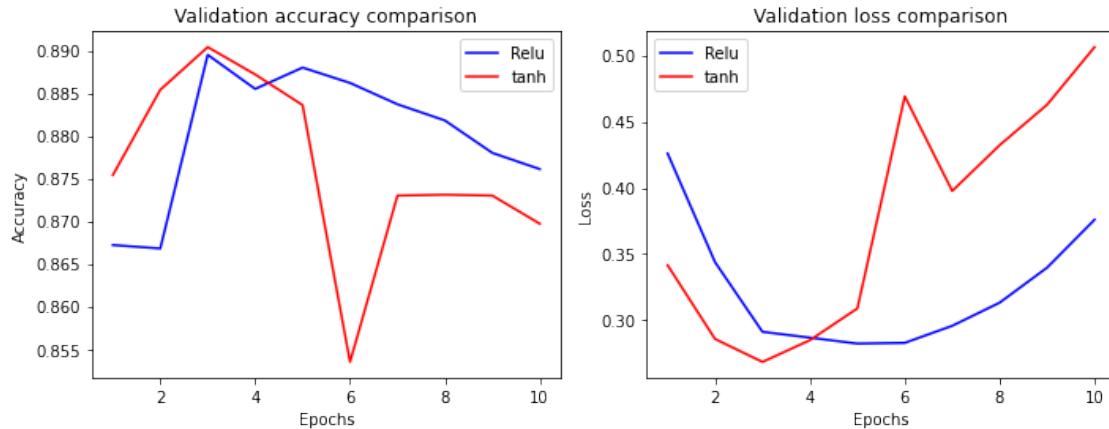
epochs = range(1, len(val_acc_tanh)+1)

plt.figure(figsize=(10,4)) # set the figure size

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="Relu")
plt.plot(epochs, val_acc_tanh, "r", label="tanh")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="Relu")
plt.plot(epochs, val_loss_tanh, "r", label="tanh")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()
```



1.4.1 Summary for Question 4:

The validation accuracy **doesnot change too much** comparing the 2 activation. As for the validation loss, **tanh activation become overfit earlier than Relu**.

Tanh activation model test accuracy is **lower than** the original Relu model.

ReLU and Tanh activation ReLU, Rectified linear unit, **output the input for positive values, and 0 otherwise**.

ReLU is the best and most advanced activation function in deep learning right now because it removes the **vanishing gradient problem** and also computationally less expensive compared to the other activation function like sigmoid and tanh.

1.5 Question 5_1: Using regularization to improve model

```
[14]: from tensorflow.keras import regularizers
```

```
[15]: # Build the model with regularization
model_regu = keras.Sequential([
    layers.Dense(16, kernel_regularizer=regularizers.l2(0.002),
    ↪activation="relu"), # use L2 regularization
    layers.Dense(16, kernel_regularizer=regularizers.l2(0.002),
    ↪activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_regu.compile(optimizer="rmsprop",
                    loss="binary_crossentropy",
                    metrics=["accuracy"])

# Train the model
```



```

history_regu = model_regu.fit(partial_x_train,
                             partial_y_train,
                             epochs=10,      # run the model for 20 epochs and each epoch
→for 512 batch
                             batch_size=512,
                             validation_data=(x_val, y_val))

```

```

Epoch 1/10
30/30 [=====] - 2s 38ms/step - loss: 0.6240 - accuracy:
0.7784 - val_loss: 0.4972 - val_accuracy: 0.8612
Epoch 2/10
30/30 [=====] - 1s 24ms/step - loss: 0.4267 - accuracy:
0.8882 - val_loss: 0.4089 - val_accuracy: 0.8779
Epoch 3/10
30/30 [=====] - 1s 24ms/step - loss: 0.3500 - accuracy:
0.9099 - val_loss: 0.3747 - val_accuracy: 0.8870
Epoch 4/10
30/30 [=====] - 1s 21ms/step - loss: 0.3145 - accuracy:
0.9195 - val_loss: 0.3635 - val_accuracy: 0.8875
Epoch 5/10
30/30 [=====] - 1s 25ms/step - loss: 0.2926 - accuracy:
0.9257 - val_loss: 0.3605 - val_accuracy: 0.8875
Epoch 6/10
30/30 [=====] - 1s 28ms/step - loss: 0.2791 - accuracy:
0.9310 - val_loss: 0.3617 - val_accuracy: 0.8843
Epoch 7/10
30/30 [=====] - 1s 27ms/step - loss: 0.2704 - accuracy:
0.9336 - val_loss: 0.3663 - val_accuracy: 0.8821
Epoch 8/10
30/30 [=====] - 1s 24ms/step - loss: 0.2613 - accuracy:
0.9367 - val_loss: 0.3662 - val_accuracy: 0.8850
Epoch 9/10
30/30 [=====] - 1s 21ms/step - loss: 0.2525 - accuracy:
0.9403 - val_loss: 0.3829 - val_accuracy: 0.8774
Epoch 10/10
30/30 [=====] - 1s 20ms/step - loss: 0.2485 - accuracy:
0.9436 - val_loss: 0.3822 - val_accuracy: 0.8764

```

```
[16]: history_dict_regu = history_regu.history
```

Comparing the validation accuracy & loss

```

[17]: val_acc_regu = history_dict_regu['val_accuracy']
      val_acc = history_dict['val_accuracy'][:10]

      val_loss_regu = history_dict_regu['val_loss']
      val_loss = history_dict['val_loss'][:10]

```

```

epochs = range(1, len(val_acc_regu)+1)


plt.figure(figsize=(10,4)) # set the figure size

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="Original")
plt.plot(epochs, val_acc_regu, "r", label="Regularization")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="Original")
plt.plot(epochs, val_loss_regu, "r", label="Regularization")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()

```



Using the regularization **doesnot** increase the validation accuracy or decrease the loss, if we zoom in the figure, the validation loss is **even increase** after regularization.

1.6 Question 5_2: Using dropout to improve model

```
[18]: # Build the model with dropout
model_drop = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dropout(0.5),
    layers.Dense(16, activation="relu"),
    layers.Dropout(0.5),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_drop.compile(optimizer="rmsprop",
                   loss="binary_crossentropy",
                   metrics=["accuracy"])

# Train the model
history_drop = model_drop.fit(partial_x_train,
                              partial_y_train,
                              epochs=10,      # run the model for 20 epochs and each epoch
                              ↪for 512 batch
                              batch_size=512,
                              validation_data=(x_val, y_val))
```

Epoch 1/10

30/30 [=====] - 2s 35ms/step - loss: 0.6375 - accuracy: 0.6392 - val_loss: 0.5536 - val_accuracy: 0.8404

Epoch 2/10

30/30 [=====] - 1s 23ms/step - loss: 0.5255 - accuracy: 0.7679 - val_loss: 0.4380 - val_accuracy: 0.8664

Epoch 3/10

30/30 [=====] - 1s 21ms/step - loss: 0.4454 - accuracy: 0.8154 - val_loss: 0.3716 - val_accuracy: 0.8689

Epoch 4/10

30/30 [=====] - 1s 24ms/step - loss: 0.3739 - accuracy: 0.8543 - val_loss: 0.3108 - val_accuracy: 0.8835

Epoch 5/10

30/30 [=====] - 1s 21ms/step - loss: 0.3233 - accuracy: 0.8778 - val_loss: 0.2867 - val_accuracy: 0.8889

Epoch 6/10

30/30 [=====] - 1s 24ms/step - loss: 0.2844 - accuracy: 0.8997 - val_loss: 0.2851 - val_accuracy: 0.8832

Epoch 7/10

30/30 [=====] - 1s 21ms/step - loss: 0.2488 - accuracy: 0.9135 - val_loss: 0.2714 - val_accuracy: 0.8939

Epoch 8/10

30/30 [=====] - 1s 23ms/step - loss: 0.2274 - accuracy: 0.9217 - val_loss: 0.2746 - val_accuracy: 0.8911

```
Epoch 9/10
30/30 [=====] - 1s 24ms/step - loss: 0.1985 - accuracy:
0.9375 - val_loss: 0.2923 - val_accuracy: 0.8925
Epoch 10/10
30/30 [=====] - 1s 24ms/step - loss: 0.1810 - accuracy:
0.9422 - val_loss: 0.2974 - val_accuracy: 0.8908
```

Validation accuracy & loss

```
[19]: history_dict_drop = history_drop.history

val_acc_drop = history_dict_drop['val_accuracy']
val_acc = history_dict['val_accuracy'][:10]

val_loss_drop = history_dict_drop['val_loss']
val_loss = history_dict['val_loss'][:10]


epochs = range(1, len(val_acc_drop)+1)

plt.figure(figsize=(10,4)) # set the figure size

plt.subplot(1, 2, 1) # row, column, index
plt.plot(epochs, val_acc, "b", label="Original")
plt.plot(epochs, val_acc_drop, "r", label="Dropout")
plt.title("Validation accuracy comparison")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, val_loss, "b", label="Original")
plt.plot(epochs, val_loss_drop, "r", label="Dropout")
plt.title("Validation loss comparison")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.tight_layout() # minimize the overlap of the two figures
plt.show()
```



Adding the dropout can slightly increase the validation accuracy and slightly lower the validation loss as well as becoming overfit in a larger epoch number.

1.6.1 Summary for Question 5:

In the current model, using **dropout method is more appropriate** to increase the validation accuracy and avoid overfit.

Using the regularization **doesnot increase the validation accuracy or decrease the loss**, if we zoom in the figure, the validation loss is **even increase** after regularization.

Dropout could be a more efficient method to lower the validation loss and increase the validation accuracy than regularization from the figure above based on this model.

1.7 Summary for this assignment:

This assignment helps us to learn the deep learning practical programming. I played with the several different model which changing the parameters in the model.

We can change **the number of hidden layers, each layers number of units, type of loss function, type of activation function** and add **regularization & dropout**(not make the model **too complex** which could **cause the overfit** happens) strategies to adjust the model.

All the **hyperparameters** changing in this practice has one goal: make the **model has better performance on the unseen test data set**.