

# Assignment4\_yli130

November 20, 2023

## 0.1 Summary for Assignment\_4

Yanxi Li, yli130@kent.edu

For each model test accuracy are as follows:

1. Model in the reference book with 600 review words: 0.886
2. Question\_1 cutoff reviews after 150 words: 0.846
3. Question\_2 retrain training sample to 160: 0.534
4. Question\_3 validate on 10,000 samples: 0.864
5. Question\_4 only top 10,000 tokens: 0.877
6. Question\_5\_1 20,000 training with embedding layer: 0.874
7. Question\_5\_2 20,000 training with pretrained word embedding: 0.876
8. Question\_5\_3 15,000 training with embedding layer: 0.867
9. Question\_5\_4 22,500 training with embedding layer: 0.88
10. Question\_5\_5 160 training with pretrained word embedding: 0.573
11. Question\_5\_6 160 training with embedding layer: 0.655
12. Question\_5\_7 2,2500 training with pretrained word embedding:0.865

From Question\_1 to Question\_4 with one-hot encoding, more reviews, more tokens, more training samples could reach higher test accuracy for the model.

The review words and tokens used in Question\_5 is 600, 20,000, respectively.

When the training data is 20,000, using the embedding layer and pretrained word embedding has the similar result, around 0.87.

As for small training dataset, like 160 training samples, although embedding test accuracy is higher than pretrained word embedding, embedding layer shows overfit soon. The training accuracy almost reach to 100% while validation accuracy only around 50%.

When the training data size is larger, for example, 2,2500, the test accuracy for embedding layer is higher than pretrained word embedding. Since the dataset has enough samples to learn, leveraging pretrained embeddings is not very helpful in this case. But for small training samples in the above illustration, pretrained embeddings worked.

## 0.2 Model in the book

### 0.2.1 Download the data

```
[1]: #!/curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
#!/tar -xzf aclImdb_v1.tar.gz
#!/rm -r aclImdb/train/unsup
```

### 0.2.2 Preparing the data

```
[2]: '''import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''
import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.2.3 Preparing integer sequence datasets

```
[3]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.2.4 A sequence model built on one-hot encoded vector sequences

```
[4]: import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
tf.one_hot (TFOpLambda)	(None, None, 20000)	0
bidirectional (Bidirectional)	(None, 64)	5128448

dropout (Dropout)	(None, 64)	0
-----		
dense (Dense)	(None, 1)	65
=====		
Total params: 5,128,513		
Trainable params: 5,128,513		
Non-trainable params: 0		
-----		

## 0.2.5 Train and test the model

```
[5]: callbacks = [
      keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm.keras",
                                     save_best_only=True)
    ]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↳callbacks=callbacks)
model = keras.models.load_model("one_hot_bidir_lstm.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```
Epoch 1/10
625/625 [=====] - 76s 116ms/step - loss: 0.5203 -
accuracy: 0.7549 - val_loss: 0.3529 - val_accuracy: 0.8786
Epoch 2/10
625/625 [=====] - 72s 115ms/step - loss: 0.3360 -
accuracy: 0.8772 - val_loss: 0.3107 - val_accuracy: 0.8756
Epoch 3/10
625/625 [=====] - 72s 115ms/step - loss: 0.2738 -
accuracy: 0.9032 - val_loss: 0.2798 - val_accuracy: 0.8922
Epoch 4/10
625/625 [=====] - 72s 115ms/step - loss: 0.2343 -
accuracy: 0.9161 - val_loss: 0.6369 - val_accuracy: 0.7568
Epoch 5/10
625/625 [=====] - 72s 115ms/step - loss: 0.2124 -
accuracy: 0.9277 - val_loss: 0.3253 - val_accuracy: 0.8788
Epoch 6/10
625/625 [=====] - 72s 115ms/step - loss: 0.1820 -
accuracy: 0.9379 - val_loss: 0.3287 - val_accuracy: 0.8904
Epoch 7/10
625/625 [=====] - 72s 115ms/step - loss: 0.1634 -
accuracy: 0.9442 - val_loss: 0.3241 - val_accuracy: 0.8696
Epoch 8/10
625/625 [=====] - 72s 115ms/step - loss: 0.1491 -
accuracy: 0.9514 - val_loss: 0.3696 - val_accuracy: 0.8784
Epoch 9/10
625/625 [=====] - 72s 115ms/step - loss: 0.1337 -
accuracy: 0.9561 - val_loss: 0.3962 - val_accuracy: 0.8832
Epoch 10/10
```

```
625/625 [=====] - 72s 115ms/step - loss: 0.1142 -  
accuracy: 0.9632 - val_loss: 0.3573 - val_accuracy: 0.8558  
782/782 [=====] - 42s 52ms/step - loss: 0.2969 -  
accuracy: 0.8859  
Test acc: 0.886
```

# Question\_1

November 20, 2023

## 0.1 Question\_1 Cutoff reviews after 150 words

### 0.1.1 Preparing the data

```
[1]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.1.2 Preparing integer sequence datasets

```
[2]: from tensorflow.keras import layers

max_length = 150
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
```

```

        lambda x, y: (text_vectorization(x), y),
        num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

```

### 0.1.3 A sequence model built on one-hot encoded vector sequences

```

[3]: import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
tf.one_hot (TFOpLambda)	(None, None, 20000)	0
bidirectional (Bidirectional)	(None, 64)	5128448
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65
Total params: 5,128,513		
Trainable params: 5,128,513		
Non-trainable params: 0		

### 0.1.4 Train and test the model

```

[4]: callbacks = [
    keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm_150words.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↪callbacks=callbacks)

```

```
model = keras.models.load_model("one_hot_bidir_lstm_150words.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```
Epoch 1/10
625/625 [=====] - 24s 32ms/step - loss: 0.5071 -
accuracy: 0.7621 - val_loss: 0.3978 - val_accuracy: 0.8272
Epoch 2/10
625/625 [=====] - 20s 31ms/step - loss: 0.3450 -
accuracy: 0.8664 - val_loss: 0.3497 - val_accuracy: 0.8576
Epoch 3/10
625/625 [=====] - 20s 31ms/step - loss: 0.2737 -
accuracy: 0.8962 - val_loss: 0.3489 - val_accuracy: 0.8652
Epoch 4/10
625/625 [=====] - 20s 31ms/step - loss: 0.2260 -
accuracy: 0.9168 - val_loss: 0.3793 - val_accuracy: 0.8592
Epoch 5/10
625/625 [=====] - 20s 31ms/step - loss: 0.1872 -
accuracy: 0.9338 - val_loss: 0.3848 - val_accuracy: 0.8578
Epoch 6/10
625/625 [=====] - 20s 31ms/step - loss: 0.1536 -
accuracy: 0.9465 - val_loss: 0.4357 - val_accuracy: 0.8556
Epoch 7/10
625/625 [=====] - 20s 31ms/step - loss: 0.1180 -
accuracy: 0.9596 - val_loss: 0.4509 - val_accuracy: 0.8110
Epoch 8/10
625/625 [=====] - 20s 31ms/step - loss: 0.0911 -
accuracy: 0.9705 - val_loss: 0.4429 - val_accuracy: 0.8320
Epoch 9/10
625/625 [=====] - 20s 31ms/step - loss: 0.0680 -
accuracy: 0.9781 - val_loss: 0.5979 - val_accuracy: 0.8522
Epoch 10/10
625/625 [=====] - 20s 31ms/step - loss: 0.0551 -
accuracy: 0.9815 - val_loss: 0.6117 - val_accuracy: 0.8420
782/782 [=====] - 12s 14ms/step - loss: 0.3837 -
accuracy: 0.8462
Test acc: 0.846
```



## Question\_2

November 20, 2023

### 0.1 Question\_2 Restrict training sample to 160

#### 0.1.1 Preparing the data

Remove the validation back to train

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''

[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos =\n'aclImdb/train/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val =\nval_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg =\n'aclImdb/train/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val =\nval_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"
```

Make a small training sample

```
[2]: '''train_q2_dir_pos = 'aclImdb/train_q2/pos'
train_dir_pos = 'aclImdb/train/pos'

for file_name in (os.listdir(train_dir_pos))[500:600]:
    train = train_dir_pos + '/' + file_name
    small_train = train_q2_dir_pos + '/' + file_name
```

```

        shutil.move(train, small_train)

train_q2_dir_neg = 'aclImdb/train_q2/neg'
train_dir_neg = 'aclImdb/train/neg'

for file_name in (os.listdir(train_dir_neg))[500:600]:
    train = train_dir_neg + '/' + file_name
    small_train = train_q2_dir_neg + '/' + file_name
    shutil.move(train, small_train)'''

```

```

[2]: "train_q2_dir_pos = 'aclImdb/train_q2/pos'\ntrain_dir_pos =
'aclImdb/train/pos'\n\nfor file_name in (os.listdir(train_dir_pos))[500:600]:\n
train = train_dir_pos + '/' + file_name\n    small_train = train_q2_dir_pos +
 '/' + file_name\n    shutil.move(train, small_train)\n    \ntrain_q2_dir_neg =
'aclImdb/train_q2/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file_name in
(os.listdir(train_dir_neg))[500:600]:\n    train = train_dir_neg + '/' +
file_name\n    small_train = train_q2_dir_neg + '/' + file_name\n
shutil.move(train, small_train)"

```

Show train, validation, test sample number

```

[3]: import os, pathlib, shutil, random
from tensorflow import keras

batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train_q2"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train_q2", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)

```

```
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 160 files belonging to 2 classes.  
Found 40 files belonging to 2 classes.  
Found 25000 files belonging to 2 classes.

### 0.1.2 Preparing integer sequence datasets

```
[4]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.3 A sequence model built on one-hot encoded vector sequences

```
[5]: import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		

input_1 (InputLayer)	[(None, None)]	0
-----		
tf.one_hot (TFOpLambda)	(None, None, 20000)	0
-----		
bidirectional (Bidirectional)	(None, 64)	5128448
-----		
dropout (Dropout)	(None, 64)	0
-----		
dense (Dense)	(None, 1)	65
=====		
Total params: 5,128,513		
Trainable params: 5,128,513		
Non-trainable params: 0		
-----		

#### 0.1.4 Train and test the model

```
[6]: callbacks = [
      keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm_smalltrain.keras",
                                     save_best_only=True)
    ]
    model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
              ↳callbacks=callbacks)
    model = keras.models.load_model("one_hot_bidir_lstm_smalltrain.keras")
    print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```
Epoch 1/10
5/5 [=====] - 5s 296ms/step - loss: 0.6950 - accuracy:
0.4500 - val_loss: 0.6929 - val_accuracy: 0.5500
Epoch 2/10
5/5 [=====] - 1s 143ms/step - loss: 0.6901 - accuracy:
0.5688 - val_loss: 0.6928 - val_accuracy: 0.5500
Epoch 3/10
5/5 [=====] - 1s 146ms/step - loss: 0.6848 - accuracy:
0.6562 - val_loss: 0.6925 - val_accuracy: 0.6000
Epoch 4/10
5/5 [=====] - 1s 144ms/step - loss: 0.6732 - accuracy:
0.8500 - val_loss: 0.6897 - val_accuracy: 0.6000
Epoch 5/10
5/5 [=====] - 1s 138ms/step - loss: 0.7405 - accuracy:
0.7875 - val_loss: 0.9221 - val_accuracy: 0.5000
Epoch 6/10
5/5 [=====] - 1s 140ms/step - loss: 0.7738 - accuracy:
0.6187 - val_loss: 0.6815 - val_accuracy: 0.5500
Epoch 7/10
5/5 [=====] - 1s 150ms/step - loss: 0.6065 - accuracy:
0.8125 - val_loss: 0.6784 - val_accuracy: 0.5750
Epoch 8/10
```

```
5/5 [=====] - 1s 144ms/step - loss: 0.5479 - accuracy:
0.8813 - val_loss: 0.6739 - val_accuracy: 0.5750
Epoch 9/10
5/5 [=====] - 1s 138ms/step - loss: 0.5641 - accuracy:
0.7875 - val_loss: 0.6730 - val_accuracy: 0.5750
Epoch 10/10
5/5 [=====] - 1s 143ms/step - loss: 0.4876 - accuracy:
0.8938 - val_loss: 0.6995 - val_accuracy: 0.5000
782/782 [=====] - 41s 52ms/step - loss: 0.6798 -
accuracy: 0.5338
Test acc: 0.534
```

## Question\_3

November 20, 2023

### 0.1 Question\_3 Validate on 10,000 samples

#### 0.1.1 Reorganize the dataset

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train_q2/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train_q2/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''

[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos =\n'aclImdb/train_q2/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val =\nval_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg =\n'aclImdb/train_q2/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val =\nval_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"

[2]: '''train_q2_dir_pos = 'aclImdb/train_q2/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(train_q2_dir_pos):
    small_train = train_q2_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(small_train, train)'''
```

```

train_q2_dir_neg = 'aclImdb/train_q2/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(train_q2_dir_neg):
    small_train = train_q2_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(small_train, train)'''

```

```

[2]: "train_q2_dir_pos = 'aclImdb/train_q2/pos'\ntrain_dir_pos =
'aclImdb/train/pos'\n\nfor file in os.listdir(train_q2_dir_pos):\n
small_train = train_q2_dir_pos + '/' + file\n    train = train_dir_pos + '/' +
file\n    shutil.move(small_train, train)\n    \ntrain_q2_dir_neg =
'aclImdb/train_q2/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file in
os.listdir(train_q2_dir_neg):\n    small_train = train_q2_dir_neg + '/' + file\n
train = train_dir_neg + '/' + file\n    shutil.move(small_train, train)"

```

### 0.1.2 Preparing the data

```

[3]: import os, pathlib, shutil, random
from tensorflow import keras

batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.4 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

```

Found 15000 files belonging to 2 classes.

Found 10000 files belonging to 2 classes.  
Found 25000 files belonging to 2 classes.

### 0.1.3 Preparing integer sequence datasets

```
[4]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.4 A sequence model built on one-hot encoded vector sequences

```
[5]: import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
tf.one_hot (TFOpLambda)	(None, None, 20000)	0



```

-----
bidirectional (Bidirectional (None, 64))          5128448
-----
dropout (Dropout)          (None, 64)            0
-----
dense (Dense)              (None, 1)              65
=====
Total params: 5,128,513
Trainable params: 5,128,513
Non-trainable params: 0
-----

```

### 0.1.5 Train and test the model

```

[6]: callbacks = [
      keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm_vali10000.keras",
                                     save_best_only=True)
    ]
    model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
              ↳callbacks=callbacks)
    model = keras.models.load_model("one_hot_bidir_lstm_vali10000.keras")
    print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```

```

Epoch 1/10
469/469 [=====] - 67s 136ms/step - loss: 0.5690 -
accuracy: 0.7110 - val_loss: 0.4378 - val_accuracy: 0.8377
Epoch 2/10
469/469 [=====] - 63s 135ms/step - loss: 0.3656 -
accuracy: 0.8679 - val_loss: 0.3033 - val_accuracy: 0.8797
Epoch 3/10
469/469 [=====] - 63s 135ms/step - loss: 0.2842 -
accuracy: 0.8983 - val_loss: 0.2830 - val_accuracy: 0.8868
Epoch 4/10
469/469 [=====] - 64s 135ms/step - loss: 0.2345 -
accuracy: 0.9210 - val_loss: 0.3388 - val_accuracy: 0.8539
Epoch 5/10
469/469 [=====] - 64s 136ms/step - loss: 0.1939 -
accuracy: 0.9363 - val_loss: 0.3432 - val_accuracy: 0.8602
Epoch 6/10
469/469 [=====] - 63s 135ms/step - loss: 0.1637 -
accuracy: 0.9448 - val_loss: 0.3528 - val_accuracy: 0.8866
Epoch 7/10
469/469 [=====] - 64s 135ms/step - loss: 0.1529 -
accuracy: 0.9502 - val_loss: 0.4134 - val_accuracy: 0.8801
Epoch 8/10
469/469 [=====] - 64s 135ms/step - loss: 0.1215 -
accuracy: 0.9621 - val_loss: 0.4512 - val_accuracy: 0.8798
Epoch 9/10

```

```
469/469 [=====] - 64s 135ms/step - loss: 0.1068 -  
accuracy: 0.9658 - val_loss: 0.4775 - val_accuracy: 0.8651  
Epoch 10/10  
469/469 [=====] - 64s 135ms/step - loss: 0.0831 -  
accuracy: 0.9743 - val_loss: 0.4113 - val_accuracy: 0.8790  
782/782 [=====] - 41s 52ms/step - loss: 0.3147 -  
accuracy: 0.8637  
Test acc: 0.864
```

# Question\_4

November 20, 2023

## 0.1 Question\_4 Only top 10,000 words

### 0.1.1 Preparing the data

```
[1]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32

'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.1.2 Preparing integer sequence datasets

```
[2]: from tensorflow.keras import layers

max_length = 600
max_tokens = 10000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.3 A sequence model built on one-hot encoded vector sequences

```
[3]: import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
tf.one_hot (TFOpLambda)	(None, None, 10000)	0
bidirectional (Bidirectional)	(None, 64)	2568448

dropout (Dropout)	(None, 64)	0
-----		
dense (Dense)	(None, 1)	65
=====		
Total params: 2,568,513		
Trainable params: 2,568,513		
Non-trainable params: 0		
-----		

#### 0.1.4 Train and test the model

```
[4]: callbacks = [
    keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm_10000words.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↳callbacks=callbacks)
model = keras.models.load_model("one_hot_bidir_lstm_10000words.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```
Epoch 1/10
625/625 [=====] - 56s 82ms/step - loss: 0.5290 -
accuracy: 0.7520 - val_loss: 0.4038 - val_accuracy: 0.8456
Epoch 2/10
625/625 [=====] - 50s 80ms/step - loss: 0.3535 -
accuracy: 0.8680 - val_loss: 0.2894 - val_accuracy: 0.8846
Epoch 3/10
625/625 [=====] - 50s 80ms/step - loss: 0.2905 -
accuracy: 0.8958 - val_loss: 0.2969 - val_accuracy: 0.8842
Epoch 4/10
625/625 [=====] - 50s 81ms/step - loss: 0.2535 -
accuracy: 0.9112 - val_loss: 0.2904 - val_accuracy: 0.8850
Epoch 5/10
625/625 [=====] - 50s 80ms/step - loss: 0.2261 -
accuracy: 0.9222 - val_loss: 0.2905 - val_accuracy: 0.8856
Epoch 6/10
625/625 [=====] - 50s 81ms/step - loss: 0.2077 -
accuracy: 0.9280 - val_loss: 0.3219 - val_accuracy: 0.8842
Epoch 7/10
625/625 [=====] - 50s 81ms/step - loss: 0.2033 -
accuracy: 0.9327 - val_loss: 0.3142 - val_accuracy: 0.8846
Epoch 8/10
625/625 [=====] - 50s 81ms/step - loss: 0.1760 -
accuracy: 0.9413 - val_loss: 0.4033 - val_accuracy: 0.8854
Epoch 9/10
625/625 [=====] - 50s 81ms/step - loss: 0.1637 -
accuracy: 0.9438 - val_loss: 0.4176 - val_accuracy: 0.8748
Epoch 10/10
```

```
625/625 [=====] - 51s 81ms/step - loss: 0.1531 -  
accuracy: 0.9505 - val_loss: 0.3691 - val_accuracy: 0.8844  
782/782 [=====] - 31s 38ms/step - loss: 0.3027 -  
accuracy: 0.8772  
Test acc: 0.877
```

# Question\_5\_1

November 20, 2023

## 0.1 Question\_5 20,000 training with embedding layer

### 0.1.1 Reorganize the dataset

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''
```

```
[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos = 'aclImdb/train/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val = val_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val = val_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"
```

### 0.1.2 Preparing the data

```
[2]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"'''
```

```

for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                     val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.1.3 Preparing integer sequence datasets

```

[3]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

```



### 0.1.4 Embedding layer model building from scratch

```
[4]: embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)
```

```
[5]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_gru.keras",
                                    save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding_1 (Embedding)	(None, None, 256)	5120000
bidirectional (Bidirectional)	(None, 64)	73984
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

625/625 [=====] - 30s 40ms/step - loss: 0.4852 - accuracy: 0.7793 - val\_loss: 0.3640 - val\_accuracy: 0.8578

Epoch 2/10

625/625 [=====] - 25s 39ms/step - loss: 0.3080 - accuracy: 0.8861 - val\_loss: 0.4552 - val\_accuracy: 0.8586

```

Epoch 3/10
625/625 [=====] - 25s 40ms/step - loss: 0.2486 -
accuracy: 0.9106 - val_loss: 0.3426 - val_accuracy: 0.8592
Epoch 4/10
625/625 [=====] - 25s 39ms/step - loss: 0.2096 -
accuracy: 0.9293 - val_loss: 0.3165 - val_accuracy: 0.8730
Epoch 5/10
625/625 [=====] - 25s 39ms/step - loss: 0.1789 -
accuracy: 0.9385 - val_loss: 0.3498 - val_accuracy: 0.8742
Epoch 6/10
625/625 [=====] - 25s 39ms/step - loss: 0.1546 -
accuracy: 0.9507 - val_loss: 0.3675 - val_accuracy: 0.8710
Epoch 7/10
625/625 [=====] - 25s 39ms/step - loss: 0.1232 -
accuracy: 0.9599 - val_loss: 0.4367 - val_accuracy: 0.8826
Epoch 8/10
625/625 [=====] - 25s 40ms/step - loss: 0.1024 -
accuracy: 0.9691 - val_loss: 0.3835 - val_accuracy: 0.8800
Epoch 9/10
625/625 [=====] - 25s 40ms/step - loss: 0.0871 -
accuracy: 0.9732 - val_loss: 0.4360 - val_accuracy: 0.8756
Epoch 10/10
625/625 [=====] - 25s 40ms/step - loss: 0.0741 -
accuracy: 0.9777 - val_loss: 0.4856 - val_accuracy: 0.8706
782/782 [=====] - 14s 17ms/step - loss: 0.3574 -
accuracy: 0.8570
Test acc: 0.857

```

### 0.1.5 Embedding layer model building with masking enabled

```

[6]: inputs = keras.Input(shape=(None,), dtype="int64")
    embedded = layers.Embedding(
        input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
    x = layers.Bidirectional(layers.LSTM(32))(embedded)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])
    model.summary()

    callbacks = [
        keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_with_masking.keras",
                                       save_best_only=True)
    ]

```

```

model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
↳callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_with_masking.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, None)]	0
embedding_2 (Embedding)	(None, None, 256)	5120000
bidirectional_1 (Bidirection	(None, 64)	73984
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

625/625 [=====] - 30s 40ms/step - loss: 0.3992 - accuracy: 0.8223 - val\_loss: 0.2784 - val\_accuracy: 0.8890

Epoch 2/10

625/625 [=====] - 23s 37ms/step - loss: 0.2313 - accuracy: 0.9112 - val\_loss: 0.2610 - val\_accuracy: 0.8904

Epoch 3/10

625/625 [=====] - 23s 37ms/step - loss: 0.1657 - accuracy: 0.9386 - val\_loss: 0.2912 - val\_accuracy: 0.8822

Epoch 4/10

625/625 [=====] - 23s 37ms/step - loss: 0.1220 - accuracy: 0.9561 - val\_loss: 0.3408 - val\_accuracy: 0.8594

Epoch 5/10

625/625 [=====] - 23s 37ms/step - loss: 0.0937 - accuracy: 0.9677 - val\_loss: 0.3690 - val\_accuracy: 0.8858

Epoch 6/10

625/625 [=====] - 23s 37ms/step - loss: 0.0704 - accuracy: 0.9760 - val\_loss: 0.3971 - val\_accuracy: 0.8808

Epoch 7/10

625/625 [=====] - 23s 37ms/step - loss: 0.0488 - accuracy: 0.9843 - val\_loss: 0.4493 - val\_accuracy: 0.8832

Epoch 8/10

625/625 [=====] - 23s 37ms/step - loss: 0.0332 - accuracy: 0.9887 - val\_loss: 0.4954 - val\_accuracy: 0.8806

```
Epoch 9/10
625/625 [=====] - 23s 37ms/step - loss: 0.0260 -
accuracy: 0.9919 - val_loss: 0.5148 - val_accuracy: 0.8812
Epoch 10/10
625/625 [=====] - 23s 37ms/step - loss: 0.0182 -
accuracy: 0.9941 - val_loss: 0.5982 - val_accuracy: 0.8746
782/782 [=====] - 13s 15ms/step - loss: 0.2978 -
accuracy: 0.8738
Test acc: 0.874
```

## Question\_5\_2

November 20, 2023

### 0.1 Question\_5\_2 20,000 training with pretrained word embedding

#### 0.1.1 Dataset preparation

```
[1]: import os, pathlib, shutil, random
from tensorflow import keras
import tensorflow as tf

batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

```
[2]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.2 Download the glove word embeddings

```
[3]: !wget http://nlp.stanford.edu/data/glove.6B.zip
!unzip -q glove.6B.zip
```

```
--2023-11-19 18:25:38-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80...
connected.
HTTP request sent, awaiting response... 302 Found
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
--2023-11-19 18:25:39-- https://nlp.stanford.edu/data/glove.6B.zip
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:443...
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]
--2023-11-19 18:25:39-- https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu
(downloads.cs.stanford.edu)|171.64.64.22|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip'

100%[=====>] 862,182,613 4.43MB/s in 2m 40s
```

2023-11-19 18:28:20 (5.13 MB/s) - 'glove.6B.zip' saved [862182613/862182613]

### 0.1.3 Parsing the GloVe word-embeddings file

```
[4]: import numpy as np
path_to_glove_file = "glove.6B.100d.txt"

embeddings_index = {}
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs

print(f"Found {len(embeddings_index)} word vectors.")
```

Found 400000 word vectors.

### 0.1.4 Preparing the GloVe word-embeddings matrix

```
[5]: embedding_dim = 100

vocabulary = text_vectorization.get_vocabulary()
word_index = dict(zip(vocabulary, range(len(vocabulary))))

embedding_matrix = np.zeros((max_tokens, embedding_dim))
for word, i in word_index.items():
    if i < max_tokens:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

embedding_layer = layers.Embedding(
    max_tokens,
    embedding_dim,
    embeddings_initializer=keras.initializers.Constant(embedding_matrix),
    trainable=False,
    mask_zero=True,
)
```

### 0.1.5 Model that uses a pretrained Embedding layer

```
[6]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = embedding_layer(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
```

```

outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("glove_embeddings_sequence_model.keras",
                                    save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        callbacks=callbacks)
model = keras.models.load_model("glove_embeddings_sequence_model.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding (Embedding)	(None, None, 100)	2000000
bidirectional (Bidirectional)	(None, 64)	34048
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 2,034,113  
 Trainable params: 34,113  
 Non-trainable params: 2,000,000

```

Epoch 1/10
625/625 [=====] - 32s 41ms/step - loss: 0.5734 -
accuracy: 0.6947 - val_loss: 0.5154 - val_accuracy: 0.7422
Epoch 2/10
625/625 [=====] - 24s 38ms/step - loss: 0.4541 -
accuracy: 0.7933 - val_loss: 0.4090 - val_accuracy: 0.8160
Epoch 3/10
625/625 [=====] - 24s 38ms/step - loss: 0.4059 -
accuracy: 0.8204 - val_loss: 0.3660 - val_accuracy: 0.8388
Epoch 4/10
625/625 [=====] - 24s 38ms/step - loss: 0.3730 -
accuracy: 0.8413 - val_loss: 0.3598 - val_accuracy: 0.8426
Epoch 5/10

```



```
625/625 [=====] - 24s 38ms/step - loss: 0.3484 -  
accuracy: 0.8530 - val_loss: 0.3452 - val_accuracy: 0.8540  
Epoch 6/10  
625/625 [=====] - 24s 38ms/step - loss: 0.3237 -  
accuracy: 0.8656 - val_loss: 0.3262 - val_accuracy: 0.8630  
Epoch 7/10  
625/625 [=====] - 24s 38ms/step - loss: 0.3092 -  
accuracy: 0.8730 - val_loss: 0.3133 - val_accuracy: 0.8682  
Epoch 8/10  
625/625 [=====] - 24s 38ms/step - loss: 0.2939 -  
accuracy: 0.8796 - val_loss: 0.3023 - val_accuracy: 0.8782  
Epoch 9/10  
625/625 [=====] - 24s 38ms/step - loss: 0.2752 -  
accuracy: 0.8877 - val_loss: 0.3071 - val_accuracy: 0.8748  
Epoch 10/10  
625/625 [=====] - 24s 38ms/step - loss: 0.2628 -  
accuracy: 0.8931 - val_loss: 0.2936 - val_accuracy: 0.8820  
782/782 [=====] - 15s 16ms/step - loss: 0.2905 -  
accuracy: 0.8761  
Test acc: 0.876
```

## Question\_5\_3

November 20, 2023

### 0.1 Question\_5\_3 15,000 training with embedding layer

#### 0.1.1 Reorganize the dataset

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''
```

```
[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos = 'aclImdb/train/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val = val_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val = val_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"
```

#### 0.1.2 Preparing the data

```
[2]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"'''
```

```

for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.4 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

```

Found 15000 files belonging to 2 classes.  
 Found 10000 files belonging to 2 classes.  
 Found 25000 files belonging to 2 classes.

### 0.1.3 Preparing integer sequence datasets

```

[3]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

```

### 0.1.4 Embedding layer model building from scratch

```
[4]: embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)
```

```
[5]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_dettrain.keras",
                                    save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_dettrain.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding_1 (Embedding)	(None, None, 256)	5120000
bidirectional (Bidirectional)	(None, 64)	73984
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

469/469 [=====] - 26s 48ms/step - loss: 0.5250 - accuracy: 0.7445 - val\_loss: 0.4734 - val\_accuracy: 0.7947

Epoch 2/10

469/469 [=====] - 21s 45ms/step - loss: 0.3327 - accuracy: 0.8736 - val\_loss: 0.3403 - val\_accuracy: 0.8611

```

Epoch 3/10
469/469 [=====] - 22s 46ms/step - loss: 0.2515 -
accuracy: 0.9119 - val_loss: 0.3305 - val_accuracy: 0.8700
Epoch 4/10
469/469 [=====] - 22s 46ms/step - loss: 0.1943 -
accuracy: 0.9303 - val_loss: 0.3693 - val_accuracy: 0.8629
Epoch 5/10
469/469 [=====] - 22s 47ms/step - loss: 0.1620 -
accuracy: 0.9430 - val_loss: 0.4676 - val_accuracy: 0.8510
Epoch 6/10
469/469 [=====] - 22s 46ms/step - loss: 0.1313 -
accuracy: 0.9581 - val_loss: 0.3755 - val_accuracy: 0.8657
Epoch 7/10
469/469 [=====] - 22s 46ms/step - loss: 0.1094 -
accuracy: 0.9653 - val_loss: 0.4294 - val_accuracy: 0.8605
Epoch 8/10
469/469 [=====] - 22s 46ms/step - loss: 0.0943 -
accuracy: 0.9723 - val_loss: 0.4606 - val_accuracy: 0.8581
Epoch 9/10
469/469 [=====] - 22s 46ms/step - loss: 0.0783 -
accuracy: 0.9771 - val_loss: 0.5311 - val_accuracy: 0.8526
Epoch 10/10
469/469 [=====] - 22s 47ms/step - loss: 0.0627 -
accuracy: 0.9809 - val_loss: 0.5050 - val_accuracy: 0.8699
782/782 [=====] - 13s 16ms/step - loss: 0.3835 -
accuracy: 0.8406
Test acc: 0.841

```

### 0.1.5 Embedding layer model building with masking enabled

```

[6]: inputs = keras.Input(shape=(None,), dtype="int64")
    embedded = layers.Embedding(
        input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
    x = layers.Bidirectional(layers.LSTM(32))(embedded)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])
    model.summary()

    callbacks = [
        keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_with_masking_detrain.
        ↳keras",
                                      save_best_only=True)
    ]

```

```

model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
↳callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_with_masking_detrain.
↳keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, None)]	0
embedding_2 (Embedding)	(None, None, 256)	5120000
bidirectional_1 (Bidirection	(None, 64)	73984
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

469/469 [=====] - 27s 47ms/step - loss: 0.4390 - accuracy: 0.7914 - val\_loss: 0.2945 - val\_accuracy: 0.8777

Epoch 2/10

469/469 [=====] - 20s 43ms/step - loss: 0.2520 - accuracy: 0.9019 - val\_loss: 0.2947 - val\_accuracy: 0.8783

Epoch 3/10

469/469 [=====] - 20s 44ms/step - loss: 0.1741 - accuracy: 0.9375 - val\_loss: 0.4056 - val\_accuracy: 0.8420

Epoch 4/10

469/469 [=====] - 20s 43ms/step - loss: 0.1231 - accuracy: 0.9568 - val\_loss: 0.3500 - val\_accuracy: 0.8822

Epoch 5/10

469/469 [=====] - 20s 43ms/step - loss: 0.0906 - accuracy: 0.9687 - val\_loss: 0.3852 - val\_accuracy: 0.8785

Epoch 6/10

469/469 [=====] - 20s 43ms/step - loss: 0.0639 - accuracy: 0.9789 - val\_loss: 0.4697 - val\_accuracy: 0.8604

Epoch 7/10

469/469 [=====] - 20s 43ms/step - loss: 0.0470 - accuracy: 0.9856 - val\_loss: 0.4737 - val\_accuracy: 0.8624

Epoch 8/10

469/469 [=====] - 21s 44ms/step - loss: 0.0338 -

```
accuracy: 0.9901 - val_loss: 0.6086 - val_accuracy: 0.8664
Epoch 9/10
469/469 [=====] - 20s 43ms/step - loss: 0.0258 -
accuracy: 0.9912 - val_loss: 0.5645 - val_accuracy: 0.8644
Epoch 10/10
469/469 [=====] - 20s 43ms/step - loss: 0.0209 -
accuracy: 0.9938 - val_loss: 0.6176 - val_accuracy: 0.8694
782/782 [=====] - 14s 15ms/step - loss: 0.3133 -
accuracy: 0.8672
Test acc: 0.867
```

## Question\_5\_4

November 20, 2023

### 0.1 Question\_5\_4 22,500 training with embedding layer

#### 0.1.1 Reorganize the dataset

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''
```

```
[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos = 'aclImdb/train/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val = val_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val = val_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"
```

#### 0.1.2 Preparing the data

```
[2]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"'''
```



```

for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.1 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

```

Found 22500 files belonging to 2 classes.

Found 2500 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.1.3 Preparing integer sequence datasets

```

[3]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

```

### 0.1.4 Embedding layer model building from scratch

```
[4]: embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)
```

```
[5]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_intrain.keras",
                                    save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_intrain.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding_1 (Embedding)	(None, None, 256)	5120000
bidirectional (Bidirectional)	(None, 64)	73984
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

704/704 [=====] - 30s 38ms/step - loss: 0.4725 - accuracy: 0.7883 - val\_loss: 0.3106 - val\_accuracy: 0.8832

Epoch 2/10

704/704 [=====] - 26s 37ms/step - loss: 0.3047 - accuracy: 0.8890 - val\_loss: 0.2784 - val\_accuracy: 0.8912

```

Epoch 3/10
704/704 [=====] - 26s 37ms/step - loss: 0.2424 -
accuracy: 0.9142 - val_loss: 0.3334 - val_accuracy: 0.8700
Epoch 4/10
704/704 [=====] - 26s 37ms/step - loss: 0.2048 -
accuracy: 0.9311 - val_loss: 0.2945 - val_accuracy: 0.8880
Epoch 5/10
704/704 [=====] - 26s 37ms/step - loss: 0.1707 -
accuracy: 0.9428 - val_loss: 0.4006 - val_accuracy: 0.8540
Epoch 6/10
704/704 [=====] - 26s 37ms/step - loss: 0.1503 -
accuracy: 0.9513 - val_loss: 0.4193 - val_accuracy: 0.8752
Epoch 7/10
704/704 [=====] - 26s 37ms/step - loss: 0.1210 -
accuracy: 0.9611 - val_loss: 0.3764 - val_accuracy: 0.8832
Epoch 8/10
704/704 [=====] - 27s 38ms/step - loss: 0.1028 -
accuracy: 0.9682 - val_loss: 0.3733 - val_accuracy: 0.8884
Epoch 9/10
704/704 [=====] - 27s 38ms/step - loss: 0.0818 -
accuracy: 0.9741 - val_loss: 0.4033 - val_accuracy: 0.8740
Epoch 10/10
704/704 [=====] - 27s 38ms/step - loss: 0.0669 -
accuracy: 0.9798 - val_loss: 0.4644 - val_accuracy: 0.8844
782/782 [=====] - 14s 17ms/step - loss: 0.3239 -
accuracy: 0.8722
Test acc: 0.872

```

### 0.1.5 Embedding layer model building with masking enabled

```

[6]: inputs = keras.Input(shape=(None,), dtype="int64")
      embedded = layers.Embedding(
          input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
      x = layers.Bidirectional(layers.LSTM(32))(embedded)
      x = layers.Dropout(0.5)(x)
      outputs = layers.Dense(1, activation="sigmoid")(x)
      model = keras.Model(inputs, outputs)
      model.compile(optimizer="rmsprop",
                    loss="binary_crossentropy",
                    metrics=["accuracy"])
      model.summary()

      callbacks = [
          keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_with_masking_intrain.
          ↳keras",
                                         save_best_only=True)
      ]

```

```

model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↳callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_with_masking_intrain.
        ↳keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, None)]	0
embedding_2 (Embedding)	(None, None, 256)	5120000
bidirectional_1 (Bidirection	(None, 64)	73984
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

704/704 [=====] - 32s 37ms/step - loss: 0.3799 - accuracy: 0.8328 - val\_loss: 0.2587 - val\_accuracy: 0.8916

Epoch 2/10

704/704 [=====] - 25s 35ms/step - loss: 0.2296 - accuracy: 0.9112 - val\_loss: 0.3617 - val\_accuracy: 0.8800

Epoch 3/10

704/704 [=====] - 25s 35ms/step - loss: 0.1714 - accuracy: 0.9366 - val\_loss: 0.2803 - val\_accuracy: 0.8980

Epoch 4/10

704/704 [=====] - 25s 35ms/step - loss: 0.1266 - accuracy: 0.9550 - val\_loss: 0.3518 - val\_accuracy: 0.8824

Epoch 5/10

704/704 [=====] - 25s 35ms/step - loss: 0.1015 - accuracy: 0.9644 - val\_loss: 0.3203 - val\_accuracy: 0.8828

Epoch 6/10

704/704 [=====] - 25s 35ms/step - loss: 0.0721 - accuracy: 0.9759 - val\_loss: 0.3958 - val\_accuracy: 0.8804

Epoch 7/10

704/704 [=====] - 25s 35ms/step - loss: 0.0536 - accuracy: 0.9821 - val\_loss: 0.4182 - val\_accuracy: 0.8800

Epoch 8/10

704/704 [=====] - 25s 35ms/step - loss: 0.0374 -

```
accuracy: 0.9872 - val_loss: 0.5612 - val_accuracy: 0.8568
Epoch 9/10
704/704 [=====] - 25s 35ms/step - loss: 0.0267 -
accuracy: 0.9911 - val_loss: 0.5322 - val_accuracy: 0.8664
Epoch 10/10
704/704 [=====] - 25s 35ms/step - loss: 0.0188 -
accuracy: 0.9939 - val_loss: 0.5649 - val_accuracy: 0.8808
782/782 [=====] - 13s 15ms/step - loss: 0.2893 -
accuracy: 0.8802
Test acc: 0.880
```

## Question\_5\_5

November 20, 2023

### 0.1 Question\_5\_5 160 training with pretrained word embedding

#### 0.1.1 Dataset preparation

```
[1]: import os, pathlib, shutil, random
from tensorflow import keras
import tensorflow as tf

batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "train_q2"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.008 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

'''train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train_q2", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
#text_only_train_ds = train_ds.map(lambda x, y: x)'''

[1]: 'train_ds = keras.utils.text_dataset_from_directory(\n    "aclImdb/train",
batch_size=batch_size\n) # change the train data directory\nval_ds =
keras.utils.text_dataset_from_directory(\n    "aclImdb/train_q2",
batch_size=batch_size\n)\ntest_ds = keras.utils.text_dataset_from_directory(\n
"aclImdb/test", batch_size=batch_size\n)\n#text_only_train_ds =
```

```
train_ds.map(lambda x, y: x)'
```

```
[2]: '''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train_q2"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train_q2", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 160 files belonging to 2 classes.

Found 40 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

```
[3]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

```
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.2 Download the glove word embeddings

```
[4]: #!/wget http://nlp.stanford.edu/data/glove.6B.zip
      #!unzip -q glove.6B.zip
```

### 0.1.3 Parsing the GloVe word-embeddings file

```
[5]: import numpy as np
path_to_glove_file = "glove.6B.100d.txt"

embeddings_index = {}
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs

print(f"Found {len(embeddings_index)} word vectors.")
```

Found 400000 word vectors.

### 0.1.4 Preparing the GloVe word-embeddings matrix

```
[6]: embedding_dim = 100

vocabulary = text_vectorization.get_vocabulary()
word_index = dict(zip(vocabulary, range(len(vocabulary))))

embedding_matrix = np.zeros((max_tokens, embedding_dim))
for word, i in word_index.items():
    if i < max_tokens:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

embedding_layer = layers.Embedding(
    max_tokens,
    embedding_dim,
    embeddings_initializer=keras.initializers.Constant(embedding_matrix),
    trainable=False,
    mask_zero=True,
)
```



### 0.1.5 Model that uses a pretrained Embedding layer

```
[7]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = embedding_layer(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("glove_embeddings_sequence_model_smalltrain.
    ↳keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
    ↳callbacks=callbacks)
model = keras.models.load_model("glove_embeddings_sequence_model_smalltrain.
    ↳keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding (Embedding)	(None, None, 100)	2000000
bidirectional (Bidirectional)	(None, 64)	34048
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 2,034,113

Trainable params: 34,113

Non-trainable params: 2,000,000

Epoch 1/10

5/5 [=====] - 8s 521ms/step - loss: 0.7336 - accuracy: 0.4938 - val\_loss: 0.6800 - val\_accuracy: 0.6500

Epoch 2/10

5/5 [=====] - 0s 61ms/step - loss: 0.6899 - accuracy: 0.5375 - val\_loss: 0.6789 - val\_accuracy: 0.6000

```
Epoch 3/10
5/5 [=====] - 0s 65ms/step - loss: 0.6779 - accuracy:
0.5750 - val_loss: 0.6731 - val_accuracy: 0.6000
Epoch 4/10
5/5 [=====] - 0s 68ms/step - loss: 0.6658 - accuracy:
0.6187 - val_loss: 0.6718 - val_accuracy: 0.6000
Epoch 5/10
5/5 [=====] - 0s 57ms/step - loss: 0.6674 - accuracy:
0.5813 - val_loss: 0.6619 - val_accuracy: 0.6250
Epoch 6/10
5/5 [=====] - 0s 64ms/step - loss: 0.6224 - accuracy:
0.6562 - val_loss: 0.6576 - val_accuracy: 0.6250
Epoch 7/10
5/5 [=====] - 0s 63ms/step - loss: 0.6200 - accuracy:
0.6500 - val_loss: 0.6552 - val_accuracy: 0.6250
Epoch 8/10
5/5 [=====] - 0s 61ms/step - loss: 0.6167 - accuracy:
0.6313 - val_loss: 0.6450 - val_accuracy: 0.6500
Epoch 9/10
5/5 [=====] - 0s 62ms/step - loss: 0.5925 - accuracy:
0.6938 - val_loss: 0.6544 - val_accuracy: 0.5500
Epoch 10/10
5/5 [=====] - 0s 55ms/step - loss: 0.5794 - accuracy:
0.7750 - val_loss: 0.6412 - val_accuracy: 0.5750
782/782 [=====] - 15s 16ms/step - loss: 0.6831 -
accuracy: 0.5732
Test acc: 0.573
```

## Question\_5\_6

November 20, 2023

### 0.1 Question\_5\_6 160 training with embedding layer

#### 0.1.1 Preparing the data

```
[1]: import os, pathlib, shutil, random
from tensorflow import keras
batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.1 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train_q2", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 160 files belonging to 2 classes.

Found 40 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

### 0.1.2 Preparing integer sequence datasets

```
[2]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.3 Embedding layer model building from scratch

```
[3]: embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)

[4]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_gru_smalltrain.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_smalltrain.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding_1 (Embedding)	(None, None, 256)	5120000
bidirectional (Bidirectional)	(None, 64)	73984
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 5,194,049

Trainable params: 5,194,049

Non-trainable params: 0

Epoch 1/10

5/5 [=====] - 4s 227ms/step - loss: 0.6955 - accuracy: 0.4688 - val\_loss: 0.6946 - val\_accuracy: 0.4500

Epoch 2/10

5/5 [=====] - 0s 93ms/step - loss: 0.6755 - accuracy: 0.6687 - val\_loss: 0.6945 - val\_accuracy: 0.5500

Epoch 3/10

5/5 [=====] - 0s 83ms/step - loss: 0.6422 - accuracy: 0.7812 - val\_loss: 0.6894 - val\_accuracy: 0.5250

Epoch 4/10

5/5 [=====] - 0s 84ms/step - loss: 0.5474 - accuracy: 0.8188 - val\_loss: 0.6898 - val\_accuracy: 0.5250

Epoch 5/10

5/5 [=====] - 0s 79ms/step - loss: 0.4489 - accuracy: 0.9375 - val\_loss: 0.6408 - val\_accuracy: 0.5750

Epoch 6/10

5/5 [=====] - 1s 94ms/step - loss: 0.2812 - accuracy: 0.9812 - val\_loss: 0.6212 - val\_accuracy: 0.6750

Epoch 7/10

5/5 [=====] - 0s 68ms/step - loss: 0.2113 - accuracy: 0.9812 - val\_loss: 0.6469 - val\_accuracy: 0.6500

Epoch 8/10

5/5 [=====] - 0s 91ms/step - loss: 0.2887 - accuracy: 0.9125 - val\_loss: 0.6402 - val\_accuracy: 0.6750

Epoch 9/10

5/5 [=====] - 0s 88ms/step - loss: 0.1542 - accuracy: 1.0000 - val\_loss: 0.6197 - val\_accuracy: 0.6750

Epoch 10/10

5/5 [=====] - 0s 62ms/step - loss: 0.1183 - accuracy: 1.0000 - val\_loss: 0.7093 - val\_accuracy: 0.6500

782/782 [=====] - 14s 17ms/step - loss: 0.6633 -  
accuracy: 0.6406  
Test acc: 0.641

#### 0.1.4 Embedding layer model building with masking enabled

```
[5]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(
    input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.
    ↳ModelCheckpoint("embeddings_bidir_gru_with_masking_smalltrain.keras",
                    save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↳callbacks=callbacks)
model = keras.models.load_model("embeddings_bidir_gru_with_masking_smalltrain.
        ↳keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, None)]	0
embedding_2 (Embedding)	(None, None, 256)	5120000
bidirectional_1 (Bidirection	(None, 64)	73984
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

Total params: 5,194,049  
Trainable params: 5,194,049  
Non-trainable params: 0

Epoch 1/10  
5/5 [=====] - 8s 503ms/step - loss: 0.6921 - accuracy: 0.5500 - val\_loss: 0.6932 - val\_accuracy: 0.5250  
Epoch 2/10  
5/5 [=====] - 0s 66ms/step - loss: 0.6627 - accuracy: 0.8625 - val\_loss: 0.6922 - val\_accuracy: 0.4750  
Epoch 3/10  
5/5 [=====] - 0s 71ms/step - loss: 0.6076 - accuracy: 0.9438 - val\_loss: 0.6864 - val\_accuracy: 0.6250  
Epoch 4/10  
5/5 [=====] - 0s 62ms/step - loss: 0.4524 - accuracy: 0.9750 - val\_loss: 0.6270 - val\_accuracy: 0.6500  
Epoch 5/10  
5/5 [=====] - 0s 62ms/step - loss: 0.3675 - accuracy: 0.8375 - val\_loss: 0.5905 - val\_accuracy: 0.6750  
Epoch 6/10  
5/5 [=====] - 0s 66ms/step - loss: 0.1127 - accuracy: 1.0000 - val\_loss: 0.6084 - val\_accuracy: 0.7000  
Epoch 7/10  
5/5 [=====] - 0s 65ms/step - loss: 0.0685 - accuracy: 1.0000 - val\_loss: 0.5916 - val\_accuracy: 0.7500  
Epoch 8/10  
5/5 [=====] - 0s 72ms/step - loss: 0.0540 - accuracy: 1.0000 - val\_loss: 0.7243 - val\_accuracy: 0.6750  
Epoch 9/10  
5/5 [=====] - 0s 69ms/step - loss: 0.0348 - accuracy: 1.0000 - val\_loss: 0.6582 - val\_accuracy: 0.7750  
Epoch 10/10  
5/5 [=====] - 0s 76ms/step - loss: 0.0317 - accuracy: 1.0000 - val\_loss: 0.6930 - val\_accuracy: 0.6750  
782/782 [=====] - 14s 15ms/step - loss: 0.6209 - accuracy: 0.6546  
Test acc: 0.655

## Question\_5\_7

November 20, 2023

### 0.1 Question\_5\_7 2,2500 training with pretrained word embedding

#### 0.1.1 Dataset preparation

```
[1]: '''import os, shutil

val_dir_pos = 'aclImdb/val/pos'
train_dir_pos = 'aclImdb/train_q2/pos'

for file in os.listdir(val_dir_pos):
    val = val_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(val, train)

val_dir_neg = 'aclImdb/val/neg'
train_dir_neg = 'aclImdb/train_q2/neg'

for file in os.listdir(val_dir_neg):
    val = val_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(val, train)'''
```

```
[1]: "import os, shutil\n\nval_dir_pos = 'aclImdb/val/pos'\ntrain_dir_pos = 'aclImdb/train_q2/pos'\n\nfor file in os.listdir(val_dir_pos):\n    val = val_dir_pos + '/' + file\n    train = train_dir_pos + '/' + file\n    shutil.move(val, train)\n\nval_dir_neg = 'aclImdb/val/neg'\ntrain_dir_neg = 'aclImdb/train_q2/neg'\n\nfor file in os.listdir(val_dir_neg):\n    val = val_dir_neg + '/' + file\n    train = train_dir_neg + '/' + file\n    shutil.move(val, train)"
```

```
[2]: '''train_q2_dir_pos = 'aclImdb/train_q2/pos'
train_dir_pos = 'aclImdb/train/pos'

for file in os.listdir(train_q2_dir_pos):
    small_train = train_q2_dir_pos + '/' + file
    train = train_dir_pos + '/' + file
    shutil.move(small_train, train)'''
```



```

train_q2_dir_neg = 'aclImdb/train_q2/neg'
train_dir_neg = 'aclImdb/train/neg'

for file in os.listdir(train_q2_dir_neg):
    small_train = train_q2_dir_neg + '/' + file
    train = train_dir_neg + '/' + file
    shutil.move(small_train, train)'''

```

```

[2]: "train_q2_dir_pos = 'aclImdb/train_q2/pos'\ntrain_dir_pos =
'aclImdb/train/pos'\n\nfor file in os.listdir(train_q2_dir_pos):\n
small_train = train_q2_dir_pos + '/' + file\n    train = train_dir_pos + '/' +
file\n    shutil.move(small_train, train)\n    \ntrain_q2_dir_neg =
'aclImdb/train_q2/neg'\ntrain_dir_neg = 'aclImdb/train/neg'\n\nfor file in
os.listdir(train_q2_dir_neg):\n    small_train = train_q2_dir_neg + '/' + file\n
train = train_dir_neg + '/' + file\n    shutil.move(small_train, train)"

```

```

[3]: import os, pathlib, shutil, random
from tensorflow import keras
import tensorflow as tf

batch_size = 32
'''base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
for category in ("neg", "pos"):
    os.makedirs(val_dir / category)
    files = (os.listdir(train_dir / category))
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.1 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname,
                    val_dir / category / fname)'''

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
) # change the train data directory
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

```

Found 22500 files belonging to 2 classes.  
Found 2500 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

```
[4]: from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

### 0.1.2 Download the glove word embeddings

```
[5]: #!/wget http://nlp.stanford.edu/data/glove.6B.zip
#!/unzip -q glove.6B.zip
```

### 0.1.3 Parsing the GloVe word-embeddings file

```
[6]: import numpy as np
path_to_glove_file = "glove.6B.100d.txt"

embeddings_index = {}
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs

print(f"Found {len(embeddings_index)} word vectors.")
```

Found 400000 word vectors.

### 0.1.4 Preparing the GloVe word-embeddings matrix

```
[7]: embedding_dim = 100

vocabulary = text_vectorization.get_vocabulary()
word_index = dict(zip(vocabulary, range(len(vocabulary))))

embedding_matrix = np.zeros((max_tokens, embedding_dim))
for word, i in word_index.items():
    if i < max_tokens:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

embedding_layer = layers.Embedding(
    max_tokens,
    embedding_dim,
    embeddings_initializer=keras.initializers.Constant(embedding_matrix),
    trainable=False,
    mask_zero=True,
)
```

### 0.1.5 Model that uses a pretrained Embedding layer

```
[8]: inputs = keras.Input(shape=(None,), dtype="int64")
embedded = embedding_layer(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("glove_embeddings_sequence_model_intrain.
    ↳keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↳callbacks=callbacks)
model = keras.models.load_model("glove_embeddings_sequence_model_intrain.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

Model: "model"

-----

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None)]	0
embedding (Embedding)	(None, None, 100)	2000000
bidirectional (Bidirectional)	(None, 64)	34048
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

Total params: 2,034,113

Trainable params: 34,113

Non-trainable params: 2,000,000

Epoch 1/10

704/704 [=====] - 33s 38ms/step - loss: 0.5665 - accuracy: 0.7031 - val\_loss: 0.4758 - val\_accuracy: 0.7672

Epoch 2/10

704/704 [=====] - 25s 35ms/step - loss: 0.4425 - accuracy: 0.7979 - val\_loss: 0.4016 - val\_accuracy: 0.8208

Epoch 3/10

704/704 [=====] - 25s 35ms/step - loss: 0.3901 - accuracy: 0.8305 - val\_loss: 0.5412 - val\_accuracy: 0.7540

Epoch 4/10

704/704 [=====] - 25s 35ms/step - loss: 0.3558 - accuracy: 0.8476 - val\_loss: 0.5107 - val\_accuracy: 0.7748

Epoch 5/10

704/704 [=====] - 25s 35ms/step - loss: 0.3316 - accuracy: 0.8612 - val\_loss: 0.3157 - val\_accuracy: 0.8676

Epoch 6/10

704/704 [=====] - 25s 35ms/step - loss: 0.3080 - accuracy: 0.8721 - val\_loss: 0.3686 - val\_accuracy: 0.8476

Epoch 7/10

704/704 [=====] - 25s 35ms/step - loss: 0.2915 - accuracy: 0.8782 - val\_loss: 0.3551 - val\_accuracy: 0.8576

Epoch 8/10

704/704 [=====] - 25s 35ms/step - loss: 0.2740 - accuracy: 0.8877 - val\_loss: 0.3211 - val\_accuracy: 0.8688

Epoch 9/10

704/704 [=====] - 25s 35ms/step - loss: 0.2616 - accuracy: 0.8952 - val\_loss: 0.3237 - val\_accuracy: 0.8624

Epoch 10/10

704/704 [=====] - 25s 35ms/step - loss: 0.2454 - accuracy: 0.9008 - val\_loss: 0.3294 - val\_accuracy: 0.8700

782/782 [=====] - 14s 16ms/step - loss: 0.3124 - accuracy: 0.8655

Test acc: 0.865